

1

Multimodal, low-quality source

Scanned pages, native PDFs, tables and figures mixed with running text -- not the clean linear English RAG tooling assumes.

2

Direct vs. translated retrieval?

Translating Bangla to English before retrieval may aid embedding alignment -- or lose information in a low-resource, morphologically rich language.

3

Which embedding model actually works?

Multilingual embedding models differ widely in training and architecture; the “multilingual paradox” shows the obvious choice can underperform.

This thesis

a localized, multimodal Bangla RAG system, evaluated empirically rather than by assumption